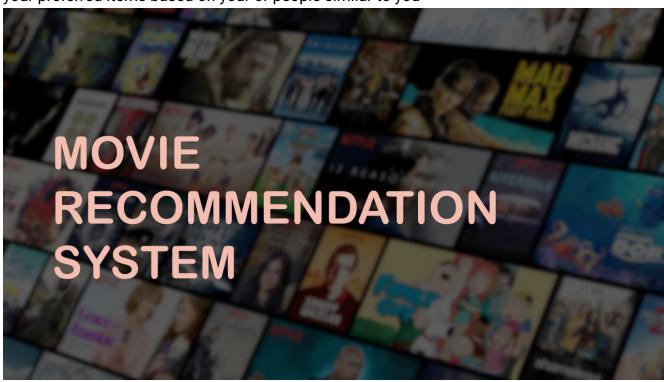
ARATHIVS

CODING_RAJA_TECHNOLOGIES-PROJECT

MOVIE RECOMMENDATION SYSTEM

A movie recommendation system is a fancy way to describe a process that tries to predict your preferred items based on your or people similar to you



IMPORTING LIBRARIES

import numpy as np
import pandas as pd

GET DATA

column_names = ['user_id', 'item_id', 'rating', 'timestamp']
df = pd.read csv("/content/u.data", sep = '\t', names = column names)

df

	user_id	item_id	rating	timestamp
0	0	50	5	881250949
1	0	172	5	881250949
2	0	133	1	881250949

×

4	186	302	3	891717742
99998	880	476	3	880175444
99999	716	204	5	879795543
100000	276	1090	1	874795795
100001	13	225	2	882399156
100002	12	203	3	879959583

100003 rows \times 4 columns

df.head()

	user_id	item_id	rating	timestamp
0	0	50	5	881250949
1	0	172	5	881250949
2	0	133	1	881250949
3	196	242	3	881250949
4	186	302	3	891717742

movie_titles=pd.read_csv("/content/Movie_Id_Titles.txt")
movie_titles

	item_id	title
0	1	Toy Story (1995)
1	2	GoldenEye (1995)
2	3	Four Rooms (1995)
3	4	Get Shorty (1995)
4	5	Copycat (1995)
1677	1678	Mat' i syn (1997)
1678	1679	B. Monkey (1998)
1679	1680	Sliding Doors (1998)
1680	1681	You So Crazy (1994)
1681	1682	Scream of Stone (Schrei aus Stein) (1991)

1682 rows \times 2 columns

df = pd.merge(df,movie_titles,on='item_id')
df.head()

	user_id	item_id	rating	timestamp	title
0	0	50	5	881250949	Star Wars (1977)
1	290	50	5	880473582	Star Wars (1977)
2	79	50	4	891271545	Star Wars (1977)
3	2	50	5	888552084	Star Wars (1977)
4	8	50	5	879362124	Star Wars (1977)

EDA

df.describe()

	user_id	item_id	rating	timestamp
count	100003.000000	100003.000000	100003.000000	1.000030e+05
mean	462.470876	425.520914	3.529864	8.835288e+08
std	266.622454	330.797791	1.125704	5.343791e+06
min	0.000000	1.000000	1.000000	8.747247e+08
25%	254.000000	175.000000	3.000000	8.794487e+08
50%	447.000000	322.000000	4.000000	8.828269e+08
75 %	682.000000	631.000000	4.000000	8.882600e+08
max	943.000000	1682.000000	5.000000	8.932866e+08

df.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 100003 entries, 0 to 100002
Data columns (total 5 columns):

#	Column	Non-Nu	ll Count	Dtype
0	user_id	100003	non-null	int64
1	item id	100003	non-null	int64
2	rating	100003	non-null	int64
3	timestamp	100003	non-null	int64
4	title	100003	non-null	object

dtypes: int64(4), object(1)

memory usage: 4.6+ MB

df.isna().sum()

```
user_id 0
item_id 0
rating 0
timestamp 0
title 0
dtype: int64
```

VISUALIZATION

```
import matplotlib.pyplot as plt
import seaborn as sns
sns.set_style('white')
%matplotlib inline
```

Lets create a ratings dataframe of average rating and number of rating

```
df.groupby('title')['rating'].mean().sort values(ascending=False).head()
    title
    They Made Me a Criminal (1939)
                                                   5.0
    Marlene Dietrich: Shadow and Light (1996)
                                                   5.0
    Saint of Fort Washington, The (1993)
                                                   5.0
    Someone Else's America (1995)
                                                   5.0
    Star Kid (1997)
                                                   5.0
    Name: rating, dtype: float64
df.groupby('title')['rating'].count().sort values(ascending=False).head()
    title
    Star Wars (1977)
                                  584
                                  509
    Contact (1997)
    Fargo (1996)
                                  508
    Return of the Jedi (1983)
                                  507
    Liar Liar (1997)
                                  485
    Name: rating, dtype: int64
ratings = pd.DataFrame(df.groupby('title')['rating'].mean())
ratings.head()
```

rating

```
title
'Til There Was You (1997) 2.333333

1-900 (1994) 2.600000

101 Dalmatians (1996) 2.908257

12 Angry Men (1957) 4.344000
```

101 (1991) 0.027000

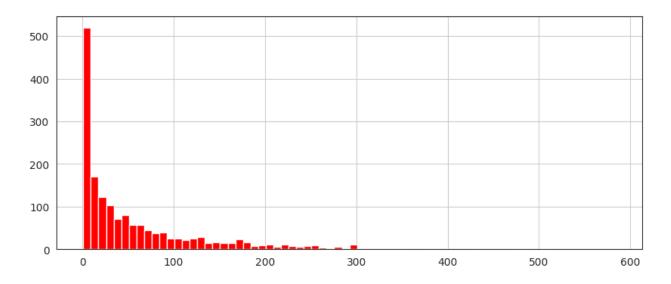
number of ratings column

ratings['num of ratings'] = pd.DataFrame(df.groupby('title')['rating'].count())
ratings.head()

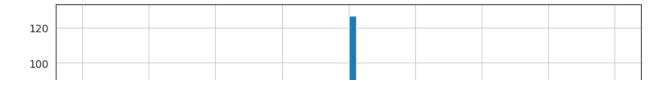
rating num of ratings

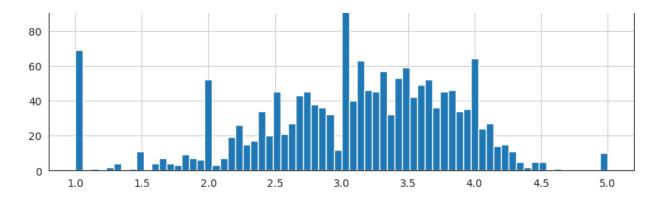
title		
'Til There Was You (1997)	2.333333	9
1-900 (1994)	2.600000	5
101 Dalmatians (1996)	2.908257	109
12 Angry Men (1957)	4.344000	125
187 (1997)	3.024390	41

plt.figure(figsize=(10,4))
ratings['num of ratings'].hist(bins=70,color='red')
plt.show()

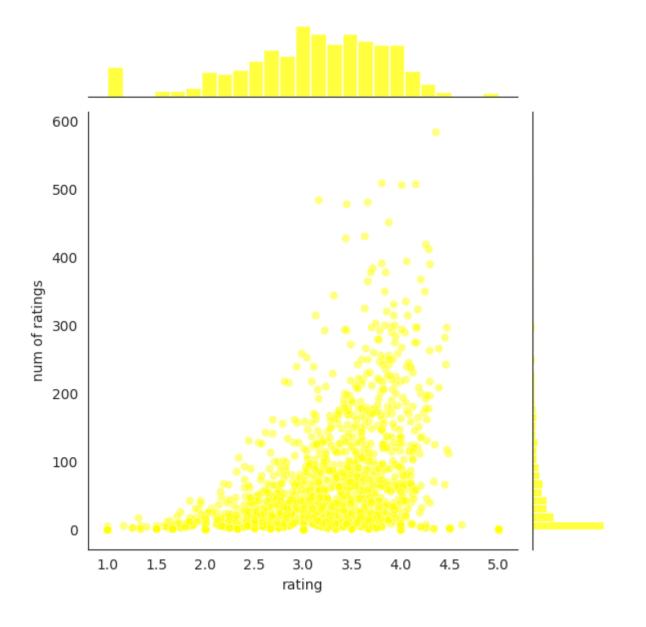


plt.figure(figsize=(10,4))
ratings['rating'].hist(bins=70)
plt.show()





sns.jointplot(x='rating',y='num of ratings', data=ratings, alpha=0.5,color='yell plt.show()



Okay! Now that we have a general idea of what the data looks like, let's move on to creating a simple recommendation system:

Recommending Similar Movies

Now let's create a matrix that has the user ids on one access and the movie title on another axis. Each cell will then consist of the rating the user gave to that movie. Note there will be a lot of NaN values, because most people have not seen most of the movies.

moviemat = df.pivot_table(index = 'user_id',columns = 'title', values = 'rating'
moviemat.head()

title user_id	'Til There Was You (1997)	1-900 (1994)	101 Dalmatians (1996)	12 Angry Men (1957)	187 (1997)	2 Days in the Valley (1996)	20,000 Leagues Under the Sea (1954)	2001: / Spac Odysse (1968
0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	Nai
1	NaN	NaN	2.0	5.0	NaN	NaN	3.0	4.0
2	NaN	NaN	NaN	NaN	NaN	NaN	NaN	Nai
3	NaN	NaN	NaN	NaN	2.0	NaN	NaN	Nai
4	NaN	NaN	NaN	NaN	NaN	NaN	NaN	Nai

 $5 \text{ rows} \times 1664 \text{ columns}$

Most rated movie:

ratings.sort values('num of ratings', ascending = False).head(10)

	rating	num of ratings
title		
Star Wars (1977)	4.359589	584
Contact (1997)	3.803536	509
Fargo (1996)	4.155512	508
Return of the Jedi (1983)	4.007890	507
Liar Liar (1997)	3.156701	485
English Patient, The (1996)	3.656965	481
Scroom (1996)	2 1/11/12	170

3CIEAIII (1330)	J.44142J	4/0
Toy Story (1995)	3.878319	452
Air Force One (1997)	3.631090	431
Independence Day (ID4) (1996)	3.438228	429

Let's choose two movies: starwars, a sci-fi movie, And Liar Liar, a comedy.

ratings.head()

rating num of ratings title 'Til There Was You (1997) 2.333333 9 1-900 (1994) 2.600000 5 **101 Dalmatians (1996)** 2.908257 109 125 12 Angry Men (1957) 4.344000 187 (1997) 41 3.024390

Now let's grab the user ratings for those two movies:

```
starwars_user_ratings = moviemat['Star Wars (1977)']
liarliar_user_ratings = moviemat['Liar Liar (1997)']
starwars_user_ratings.head()
```

```
user_id

0    5.0

1    5.0

2    5.0

3    NaN

4    5.0

Name: Star Wars (1977), dtype: float64
```

We can then use corrwith() method to get correlations between two pandas series:

```
similar_to_starwars = moviemat.corrwith(starwars_user_ratings)
similar_to_liarliar = moviemat.corrwith(liarliar_user_ratings)

/usr/local/lib/python3.10/dist-packages/numpy/lib/function_base.py:2821: Ri
    c = cov(x, y, rowvar, dtype=dtype)
/usr/local/lib/python3.10/dist-packages/numpy/lib/function_base.py:2680: Ri
    c *= np.true_divide(1, fact)
/usr/local/lib/python3.10/dist-packages/numpy/lib/function_base.py:2821: Ri
    c = cov(x, y, rowvar, dtype=dtype)
```

```
/usr/local/lib/python3.10/dist-packages/numpy/lib/function_base.py:2680: Ru
    c *= np.true_divide(1, fact)
```

Let's clean this by removing NaN values and using a DataFrame instead of a series:

```
corr_starwars = pd.DataFrame(similar_to_starwars, columns = ['Correlation'])
corr_starwars.dropna(inplace=True)
corr_starwars.head()
```

Correlation

title	
'Til There Was You (1997)	0.872872
1-900 (1994)	-0.645497
101 Dalmatians (1996)	0.211132
12 Angry Men (1957)	0.184289
187 (1997)	0.027398

Now if we sort the dataframe by correlation, we should get the most similar movies, however note that we get some results that don't really make sense. This is because there are a lot of movies only watched once by users who also watched star wars (it was the most popular movie).

corr starwars.sort values('Correlation',ascending=False).head(10)

Correlation

title	
Hollow Reed (1996)	1.0
Commandments (1997)	1.0
Cosi (1996)	1.0
No Escape (1994)	1.0
Stripes (1981)	1.0
Star Wars (1977)	1.0
Man of the Year (1995)	1.0
Beans of Egypt, Maine, The (1994)	1.0
Old Lady Who Walked in the Sea, The (Vieille qui marchait dans	1 0

la mer, La) (1991)	1.0
Outlaw, The (1943)	1.0

Let's fix this by filtering out movies that have less than 100 reviews (this value was chosen based off the histogram from earlier).

corr_starwars = corr_starwars.join(ratings['num of ratings'])
corr starwars.head()

	Correlation	num of ratings
title		
'Til There Was You (1997)	0.872872	9
1-900 (1994)	-0.645497	5
101 Dalmatians (1996)	0.211132	109
12 Angry Men (1957)	0.184289	125
187 (1997)	0.027398	41

Now sort the values and notice how the titles make a lot more sense:

corr_starwars[corr_starwars['num of ratings']>100].sort_values('Correlation', as

	Correlation	num of ratings
titl	Le	
Star Wars (1977)	1.000000	584
Empire Strikes Back, The (1980)	0.748353	368
Return of the Jedi (1983)	0.672556	507
Raiders of the Lost Ark (1981)	0.536117	420
Austin Powers: International Man of Mystery	∩ २ ७७/२२	130

Now the same for the comedy Liar Liar:

corr_liarliar = pd.DataFrame(similar_to_liarliar, columns = ['Correlation'])
corr_liarliar.dropna(inplace = True)
corr_liarliar.head()

Correlation

title
'Til There Was You (1997) 0.118913

101 Dalmatians (1996)	0.469765
12 Angry Men (1957)	0.066272
187 (1997)	0.175145
2 Days in the Valley (1996)	0.040739

corr_liarliar = corr_liarliar.join(ratings['num of ratings'])
corr_liarliar[corr_liarliar['num of ratings']>100].sort_values('Correlation', as

Correlation num of ratings

title

Liar Liar (1997)	1.000000	485
Batman Forever (1995)	0.516968	114
Mask, The (1994)	0.484650	129
Down Periscope (1996)	0.472681	101
Con Air (1997)	0.469828	137

Colab paid products - Cancel contracts here